

1 Appendix 1




Fig. 1 Each pair of plots represents one run that did not need human help. The line show the metrics of success in the validation stages.

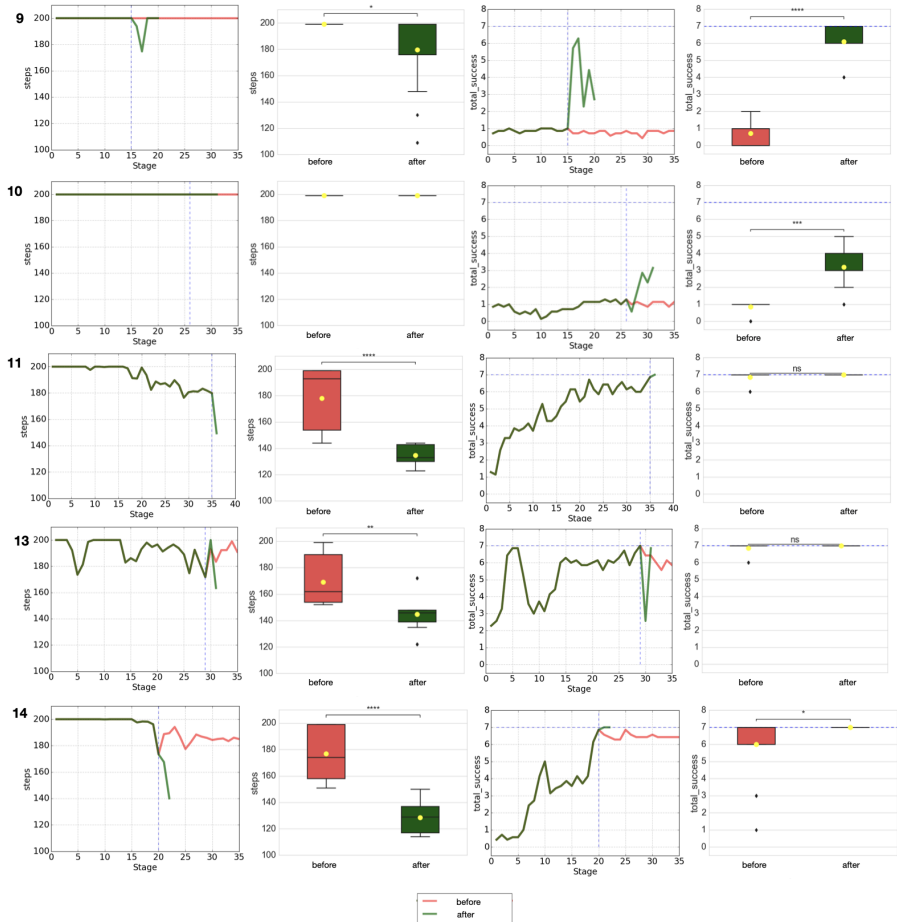
2 *Hu-bot cooperation*

Fig. 2 Each row (made of four plots) represents one run that improved after human help. before=experiment before human help and after=experiment after human help. The line plots show the metrics of success in the validation stages throughout the learning processes before and after human help: note that naturally, the lines overlap until the stage when the human help starts and the human help lines can be shorter. The boxplots compare the tests of the best controllers found by each experiment. The number of the run is displayed at the left of the plots. The horizontal dotted lines mark the maximum number of packages, while the vertical ones mark the best stage of the validation of the run before human help.